



Tasking in OpenMP

Alejandro Duran

Barcelona Supercomputing Center

Outline

- 1 Why task parallelism?
- 2 The OpenMP tasking model
 - Creating tasks
 - Data scoping
 - Synchronizing tasks
 - Execution model
- 3 Pitfalls & Performance issues
- 4 Conclusions



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Why task parallelism?

List traversal

Example

```
void traverse_list ( List l )
{
    Element e;

    #pragma omp parallel private(e)
        for ( e = l->first; e ; e = e->next )
            #pragma omp single nowait
                process(e);
}
```

Without tasks

- Awkward
- Very poor performance
- Not composable



Why task parallelism?

Tree traversal

Example

```
void traverse (Tree *tree)
{
#pragma omp parallel sections
{
#pragma omp section
    if ( tree->left )
        traverse(tree->left);
#pragma omp section
    if ( tree->right )
        traverse(tree->right);
}

process(tree);
}
```

Without tasks

- Too many parallel regions
 - Extra overheads
 - Extra synchronizations
 - Not always well supported



Task parallelism

- Better solution for those problems
- Main addition to OpenMP 3.0^a
- Allows to parallelize irregular problems
 - unbounded loops
 - recursive algorithms
 - producer/consumer schemes
 - ...

^aAyguadé et al., The Design of OpenMP Tasks, IEEE TPDS March 2009



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



What is an OpenMP task?

- Tasks are work units which execution **may** be deferred
 - they can also be executed immediately!
- Tasks are composed of:
 - **code** to execute
 - **data** environment
 - Initialized at creation time
 - **internal control variables** (ICVs)



Task directive

```
#pragma omp task [ clauses ]
    structured block
```

- Each encountering **thread** creates a task
 - Packages code and data environment
- Highly **composable**. Can be nested
 - inside parallel regions
 - inside other tasks
 - inside worksharing



List traversal with tasks

Example

```
void traverse_list ( List l )
{
Element e;
for ( e = l->first; e ; e = e->next )
    #pragma omp task
        process(e);
}
```



List traversal with tasks

Example

```
void traverse_list ( List l )
{
Element e;
for ( e = l->first; e ; e = e->next )
    #pragma omp task
        process(e); ←
}
```

What is the scope of e?



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Task data scoping

Data scoping clauses

- `shared(list)`
- `private(list)`
- `firstprivate(list)`
 - data is captured at creation
- `default(shared|none)`



Task data scoping

When there are no clauses ...

If no clause

- **Implicit rules apply**
 - e.g., global variables are shared
- Otherwise...
 - `firstprivate`
 - `shared` attributed is lexically inherited



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a =
            b =
            c =
            d =
            e =
        }}}
```



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a = shared
            b =
            c =
            d =
            e =
        }}}
```



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a = shared
            b = firstprivate
            c =
            d =
            e =

        }}}
```



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a = shared
            b = firstprivate
            c = shared
            d =
            e =

        }}}
```



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a = shared
            b = firstprivate
            c = shared
            d = firstprivate
            e =
        }
    }
}
```



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a = shared
            b = firstprivate
            c = shared
            d = firstprivate
            e = private
        }
    }
}
```



Task data scoping

In practice...

Example

```
int a;
void foo() {
    int b,c;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d;
        #pragma omp task
        {
            int e;

            a = shared
            b = firstprivate
            c = shared
            d = firstprivate
            e = private
        }
    }
}
```

Tip

default(*none*) is your friend

- Use it if you do not see it clear



List traversal

Example

```
void traverse_list ( List l )
{
Element e;
for ( e = l->first; e ; e = e->next )
    #pragma omp task
        process(e); ← e is firstprivate
}
```



List traversal

Example

```
void traverse_list ( List l )
{
Element e;
for ( e = l->first; e ; e = e->next )
    #pragma omp task
        process(e);
}
```

how we can guarantee here that the traversal is finished?



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- **Synchronizing tasks**
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Task synchronization

- Barriers (implicit or explicit)
 - All tasks created by any thread of the current team are guaranteed to be completed at barrier exit
- Task barrier

```
#pragma omp taskwait
```

 - Encountering task suspends until **child** tasks complete
 - Only **direct childs** not descendants!



List traversal

Example

```
void traverse_list ( List l )
{
Element e;
for ( e = l->first; e ; e = e->next )
    #pragma omp task
        process(e);

```

```
#pragma omp taskwait
```

← All tasks guaranteed to be completed here



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Task execution model

- Task are executed by a thread of the **team** that generated it
 - Can be executed **immediately** by the same thread that creates it
- Parallel regions in 3.0 create tasks!
 - One **implicit** task is created for each thread
 - So all task-concepts have sense inside the parallel region
- Threads can **suspend** the execution of a task and **start/resume** another



List traversal

Example

List 1

```
#pragma omp parallel  
traverse_list(l);
```



List traversal

Example

List l

```
#pragma omp parallel  
traverse_list(l);
```

Careful!

Multiple traversals of the same list



List traversal

Single traversal

Example

List 1

```
#pragma omp parallel
#pragma omp single
    traverse_list(1);
```

Single traversal

- One thread enters **single** and creates all tasks
- All the team cooperates executing them



List traversal

Multiple traversals

Example

```
List l[N]
```

```
#pragma omp parallel
#pragma omp for
for (i = 0; i < N; i++)
    traverse_list(l[i]);
```

Multiple traversals

- Multiple threads create tasks
- All the team cooperates executing them



Task scheduling

How it works?

- Tasks are **tied** by default
 - Tied tasks are executed always by the same thread
 - Tied tasks have scheduling restrictions
 - Deterministic scheduling points (creation, synchronization, ...)
 - Another constraint to avoid deadlock problems
 - Tied tasks may run into performance problems
- Programmer can use **untied** clause to lift all restrictions
 - **Note:** Mix **very carefully** with threadprivate, critical and thread-ids



And last...

The IF clause

- If the expression of a `if` clause evaluates to `false`
 - The encountering task is suspended
 - The new task is **executed immediately**
 - with its own data environment
 - different task with respect to synchronization
 - The parent task resumes when the task finishes
 - Allows implementations to **optimize** task creation



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)

    {
        state[j] = i;
        if (ok(j+1,state)) {
            search(n,j+1,state);
        }
    }
}
```



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
        #pragma omp task
    {
        state[j] = i;
        if (ok(j+1,state)) {
            search(n,j+1,state);
        }
    }
}
```



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
        #pragma omp task
    {
        state[j] = i;
        if (ok(j+1,state)) {
            search(n,j+1,state);
        }
    }
}
```

Data scoping

Because it's an **orphaned** task all variables are **firstprivate**



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
        #pragma omp task
    {
        state[j] = i;
        if (ok(j+1,state)) {
            search(n,j+1,state);
        }
    }
}
```

Data scoping

Because it's an **orphaned task** all variables are **firstprivate**

State is not captured

Just the pointer is captured
not the pointed data



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
        #pragma omp task
    {
        state[j] = i;
        if (ok(j+1,state)) {
            search(n,j+1,state);
        }
    }
}
```

Pitfall #1

Incorrectly capturing
pointed data



Pitfall #1

Incorrectly capturing pointed data

Problem

`firstprivate` does not allow to capture data through pointers

Solutions

- ① Capture it manually
- ② Copy it to an array and capture the array with `firstprivate`



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state,state,sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }
}
```



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state,state,sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }
}
```

Caution!

Will `new_state` still be valid by the time `memcpy` is executed?



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state,state,sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }
}
```

Pitfall #2

Data can go out of scope!



Pitfall #2

Out-of-scope data

Problem

Stack-allocated parent data can become invalid before being used by child tasks

- Only if not captured with `firstprivate`

Solutions

- 1 Use `firstprivate` when possible
- 2 Allocate it in the heap
 - Not always easy (we also need to free it)
- 3 Put additional synchronizations
 - May reduce the available parallelism



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }

# pragma omp taskwait
}
```



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , c
        solutions++ ←
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }

# pragma omp taskwait
}
```

Shared variable needs protected access



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }

#pragma omp taskwait
}
```

Solutions

- Use `omp critical`
- Use `omp atomic`
- Use a reduction operation
 - Not available for 3.0
 - Can be work out manually



Reductions for tasks

Example

```
int solutions=0;
int mysolutions=0; ← Use a separate counter for each thread
#pragma omp threadprivate(n)

void start_search ()
{
    #pragma omp parallel
    {
        #pragma omp single
        {
            bool initial_state[n];
            search(n,0,initial_state );
        }
        #pragma omp critical
        solutions += mysolutions; ← Accumulate them at the end
    }
}
```



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state))←
            search(n,j+1,new_state);
    }

# pragma omp taskwait
}
```

Pruning mechanism potentially
introduces imbalance in the tree



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task untied
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)){
            search(n,j+1,new_state);
        }
    }

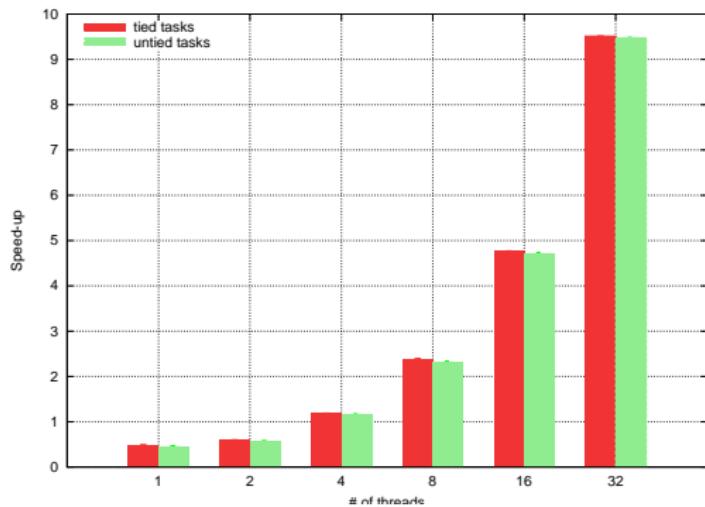
#pragma omp taskwait
}
```

Untied clause

- Allows the implementation to easier load balance



Benefit of untied?



with Intel's icc v11.0

- Don't expect much today.
- But, as implementations are optimized differences may arise



Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        solutions++ ←
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
#pragma omp task untied
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }

#pragma omp taskwait
}
```

Because of **untied** this is **not safe!**



Pitfall #3

Unsafe use of untied tasks

Problem

Because tasks can migrate between threads at any point
thread-centric constructs can yield unexpected results

Remember

When using **untied** tasks avoid:

- Threadprivate variables
- Any thread-id uses

And be very careful with:

- Critical regions (and locks)

Simple solution

Create a task tied region with **#pragma omp task if(0)**

Search problem

Example

```
void search (int n, int j, bool *state)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        #pragma omp task if(0)
        solutions++ ←
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
    #pragma omp task untied
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state);
        }
    }

    #pragma omp taskwait
}
```

Now this statement is tied and safe



Task granularity

Granularity is a key performance factor

- Tasks tend to be fine-grained
- Try to “group” tasks together
 - Use if clause or manual transformations



Using the if clause

Example

```
void search (int n, int j, bool *state, int depth)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        #pragma omp task if(0)
        mysolutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
    #pragma omp task untied if(depth < MAX_DEPTH)
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state,state,sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            search(n,j+1,new_state,depth+1);
        }
    }
    #pragma omp taskwait
}
```



Using an if statement

Example

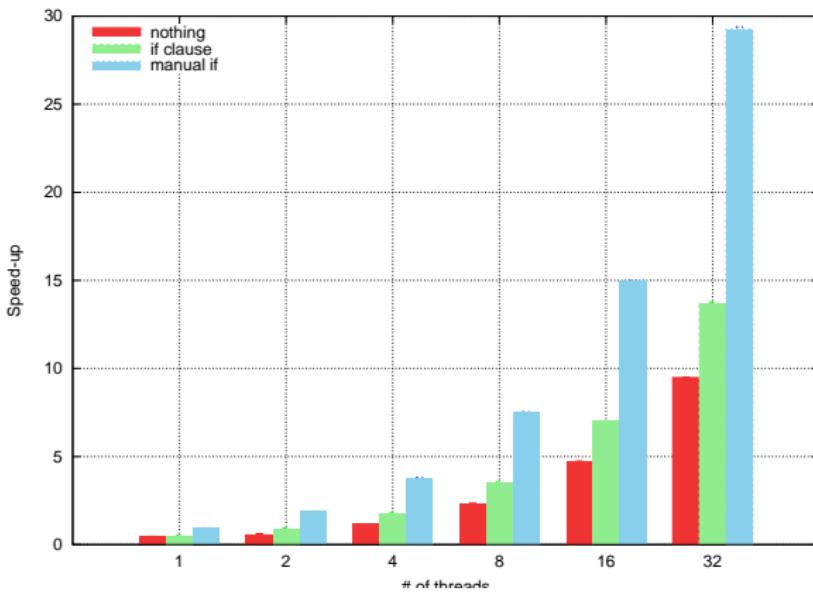
```
void search (int n, int j, bool *state, int depth)
{
    int i,res;

    if (n == j) {
        /* good solution , count it */
        #pragma omp task if(0)
        mysolutions++;
        return;
    }

    /* try each possible solution*/
    for (i = 0; i < n; i++)
    #pragma omp task untied
    {
        bool *new_state = alloca(sizeof(bool)*n);
        memcpy(new_state , state , sizeof(bool)*n);
        new_state[j] = i;
        if (ok(j+1,new_state)) {
            if (depth < MAX_DEPTH)
                search(n,j+1,new_state,depth+1);
            else
                search_serial(n,j+1,new_state);
        }
    }
    #pragma omp taskwait
}
```



If clause vs If statement



with Intel's icc v11.0

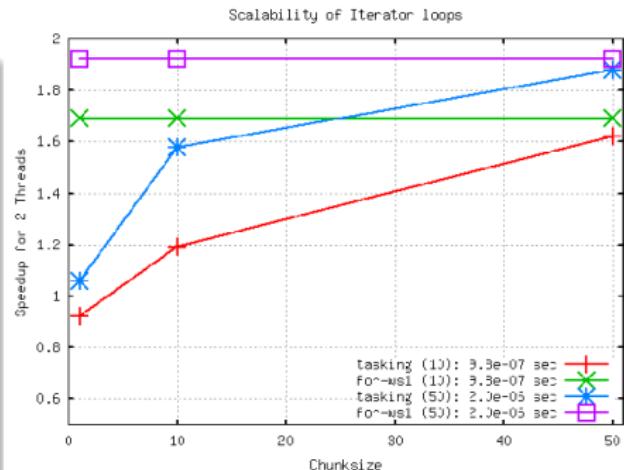
- **If clause** reduces overheads
 - without modifying the code
- but if granularity is very small is not enough



Don't abuse tasks

Tasks are nice but...

- They are not the answer to everything
 - They are more costly than other OpenMP mechanisms
- Use other OpenMP constructs when appropriate
 - Particularly for/do worksharing and sections



Courtesy of Christian Terboven



Outline

1 Why task parallelism?

2 The OpenMP tasking model

- Creating tasks
- Data scoping
- Synchronizing tasks
- Execution model

3 Pitfalls & Performance issues

4 Conclusions



Summary

Tasks in 3.0

- `#pragma omp task [clauses]` creates a task
 - `shared,firstprivate,private` data clauses
 - `firstprivate` is usually the default (but shared inherited)
 - `untied` allows tasks to move between threads
 - `if` allows to dynamically control task creation
- `#pragma omp taskwait` waits for children completion



Summary

Pitfalls & tips

- ① Use `default(None)` if unsure of data scoping
- ② Careful when using `firstprivate` on pointers
- ③ Careful with Out-of-scope data
- ④ Use `untied` tasks carefully
- ⑤ Control granularity
- ⑥ Do not abuse of tasks



Tasks after 3.0

In the works

- Support for task reductions
- Task dependences
- Scheduling hints



The End

Thanks for your attention!

